



Congestion Control

Congestion, TCP Tahoe, slow start


Congestion Control Motivation



Congestion Control Motivation

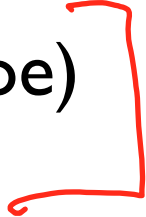


Congestion Control Motivation



Congestion control: control data rate so it does not congest network, improves overall performance

TCP History

- 1974: 3-way handshake
 - 1978: TCP and IP split into TCP/IP
 - 1983: January 1, ARPAnet switches to TCP/IP
 - 1986: Internet begins to suffer congestion collapse
 - 1987-8: Van Jacobson fixes TCP, publishes seminal TCP paper (Tahoe)
 - 1990: Fast recovery and fast retransmit added (Reno)
- 

Three Questions


- When should you send new data?
- When should you send data retransmissions?
- When should you send acknowledgments?

TCP Pre-Tahoe

- Endpoint has the flow control window size
- On connection establishment, send a full window of packets
- Start a retransmit timer for each packet
- Problem: what if window is much larger than what network can support?

window 30KB

TCP in 1986



Three Improvements

- Congestion window
 - Timeout estimation
 - Self-clocking
- Future*

Three Improvements

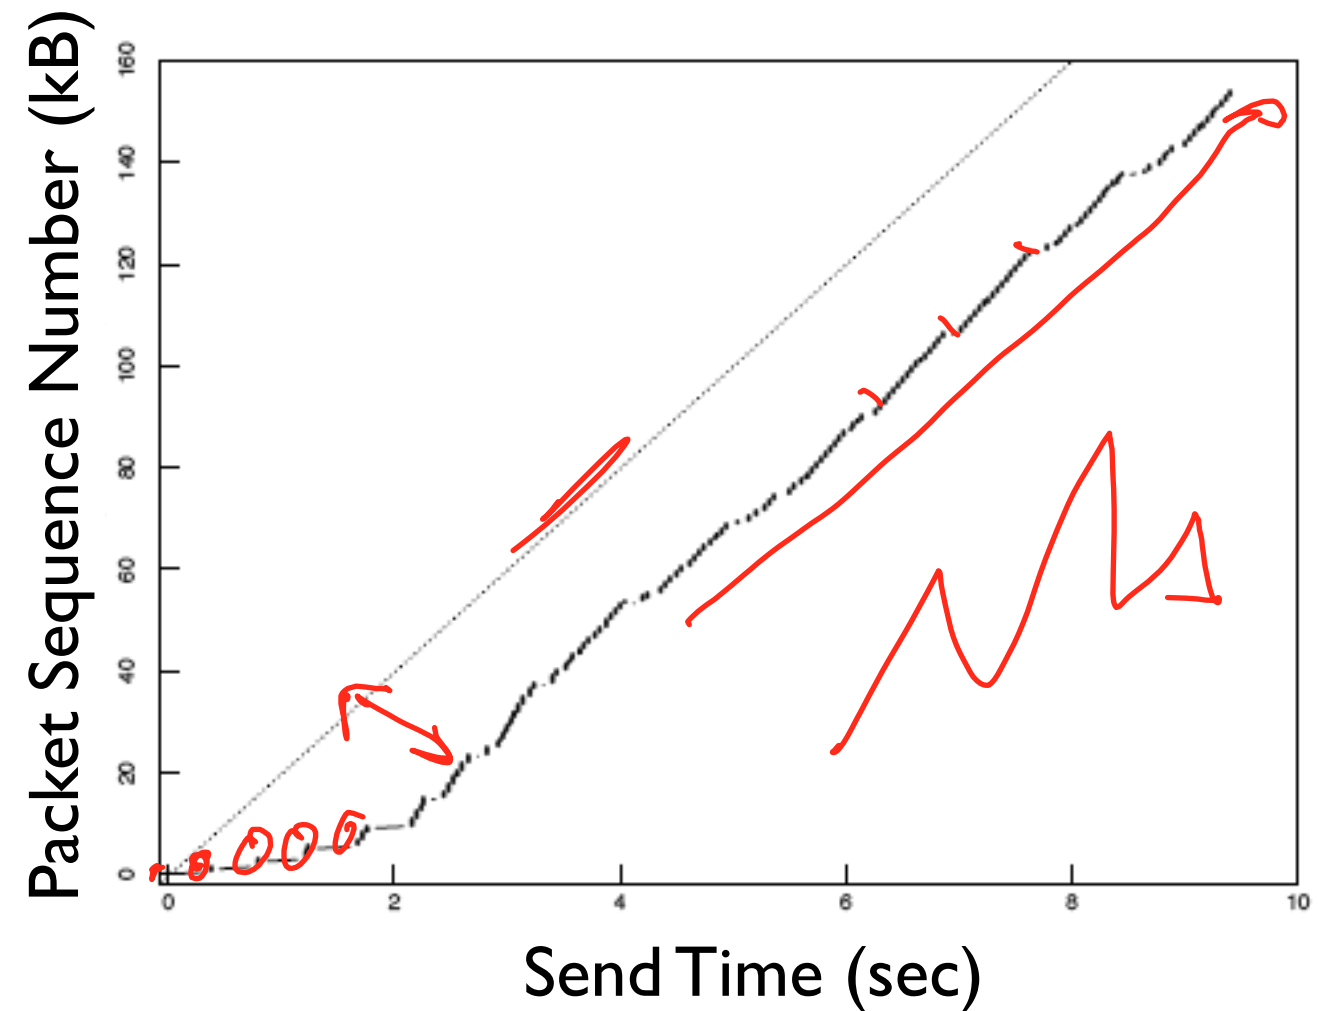
- Congestion window
- Timeout estimation
- Self-clocking

Congestion Window (TCP Tahoe)

- Flow control window is only about endpoint
- Have TCP estimate a *congestion window* for the network
- Sender window = $\min(\text{flow window}, \text{congestion window})$
- Separate congestion control into two states
 - ▶ Slow start: on connection startup or packet timeout
 - ▶ Congestion avoidance: steady operation



Slow Start Benefits

- Slow start
 - ▶ Window starts at Maximum Segment Size (MSS)
 - ▶ Increase window by MSS for each acknowledged packet
- Exponentially grow congestion window to sense network capacity
- “Slow” compared to prior approach



Congestion Avoidance

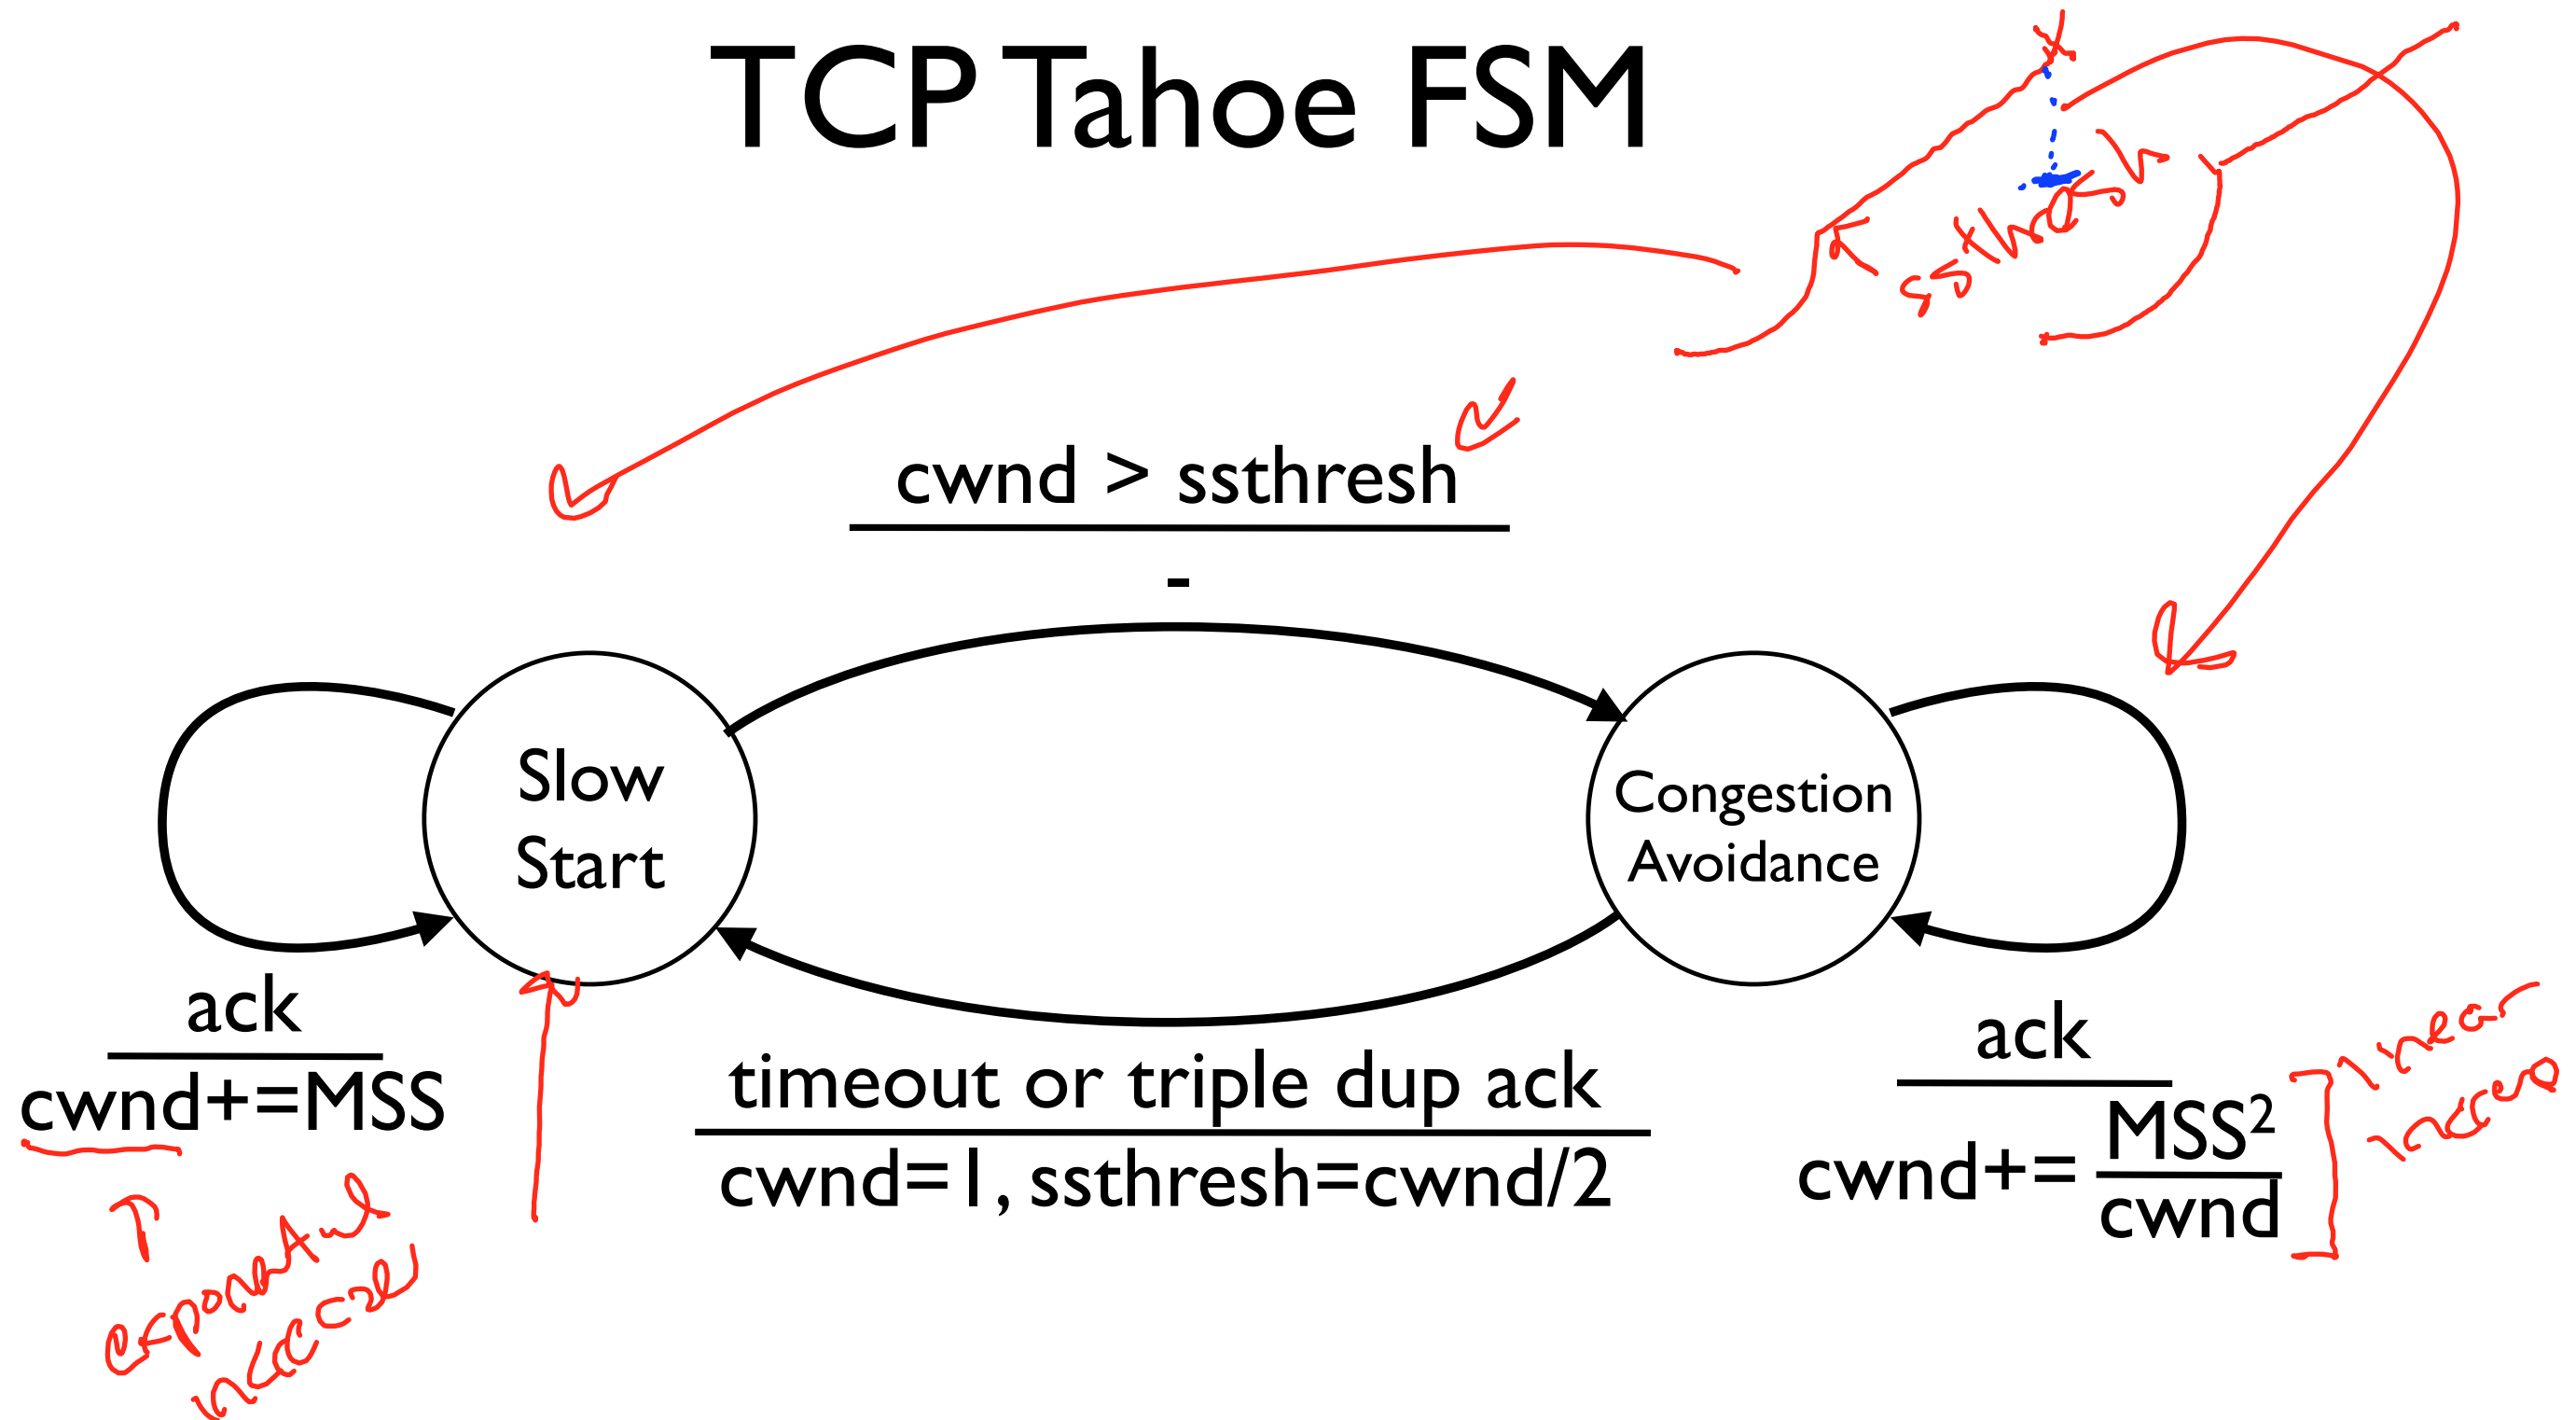
- Slow start
 - ▶ Increase congestion window by MSS for each acknowledgment
 - ▶ Exponential increase
- Congestion avoidance
 - ▶ Increase by $MSS^2 / \text{congestion window}$ for each acknowledgment
 - ▶ Behavior: increase by MSS each round trip time
 - ▶ Linear increase



State Transitions

- Two goals
 - ▶ Use slow start to quickly find network capacity
 - ▶ When close to capacity, use congestion avoidance to very carefully probe
- Three signals
 - ▶ Increasing acknowledgments: transfer is going well
 - ▶ Duplicate acknowledgments: something was lost/delayed
 - ▶ Timeout: something is very wrong

TCP Tahoe FSM



TCP Tahoe Walkthrough

