

Oh End-to-End, Where Art Thou?

Jana Iyengar

Franklin & Marshall College

Acknowledgments

- Syed Obaid Amin (F&M, Yale)
- Fitz Nowlan (Yale)
- Nabin Tiwari (F&M)
- Jeff Wise (F&M)
- Bryan Ford (Yale)

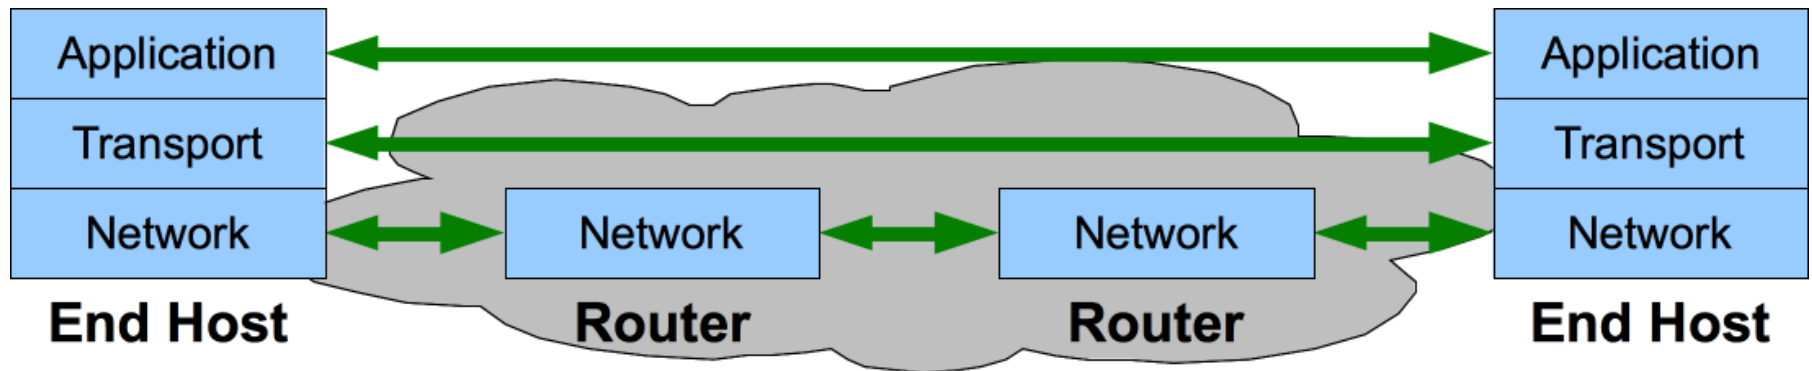
What is the Internet?

Why is it so different from other networks we've seen before?

The End-to-End Principle

Only end hosts

- see past a packet's IP header (**Generality**)
- maintain “hard state” (**Fate Sharing**)



Rise of the Middle

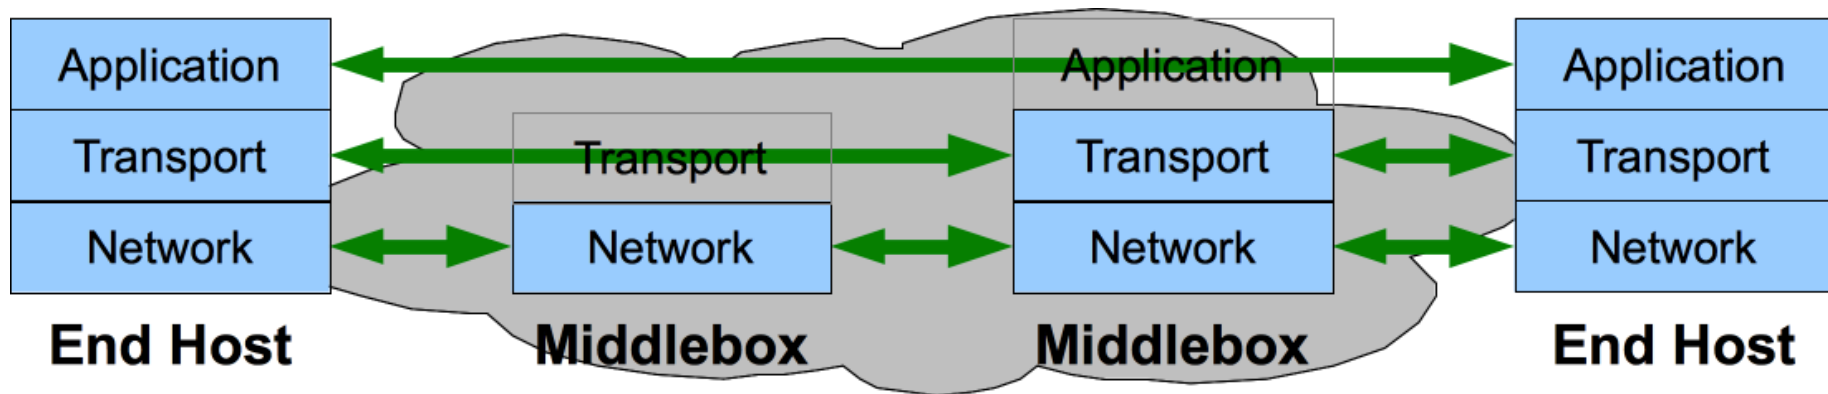
- NATs: IP address scarcity
- Firewalls: protection and policy
- Traffic Shapers: manage bandwidth and delay
- PEPs: optimize performance

All solve important problems for operators

Eroding Transport End-to-Endness

Middleboxes **need to** interact with Transport

- NATs: port number, checksum
- Others: most TCP header fields, state machine



Eroding Transport End-to-Endness

We lose:

- Generality: only TCP, UDP
- Fate Sharing: hard state inside network

What is “Reachability” now?

Middleboxes are new architectural control points ...

(control what the Internet *can* do)

... but they are accidental control points!

Accidental Control Points

RFC 1631 (NAT):

“The two most compelling problems facing the IP Internet are IP address depletion and scaling in routing. [...] Until the long-term solutions are ready an easy way to hold down the demand for IP addresses is through address reuse.”

“NAT has several negative characteristics that make it inappropriate as a long term solution, and may make it inappropriate even as a short term solution.”

A Flaky Internet!

- Packet blackholes
 - Will a particular packet will get through?
 - On what paths?
- Cheap boxes are common
 - And buggy!
 - Almost all our traffic goes through these boxes!

Do we need new Transports?

Isn't TCP enough?

Consider e2e multipath

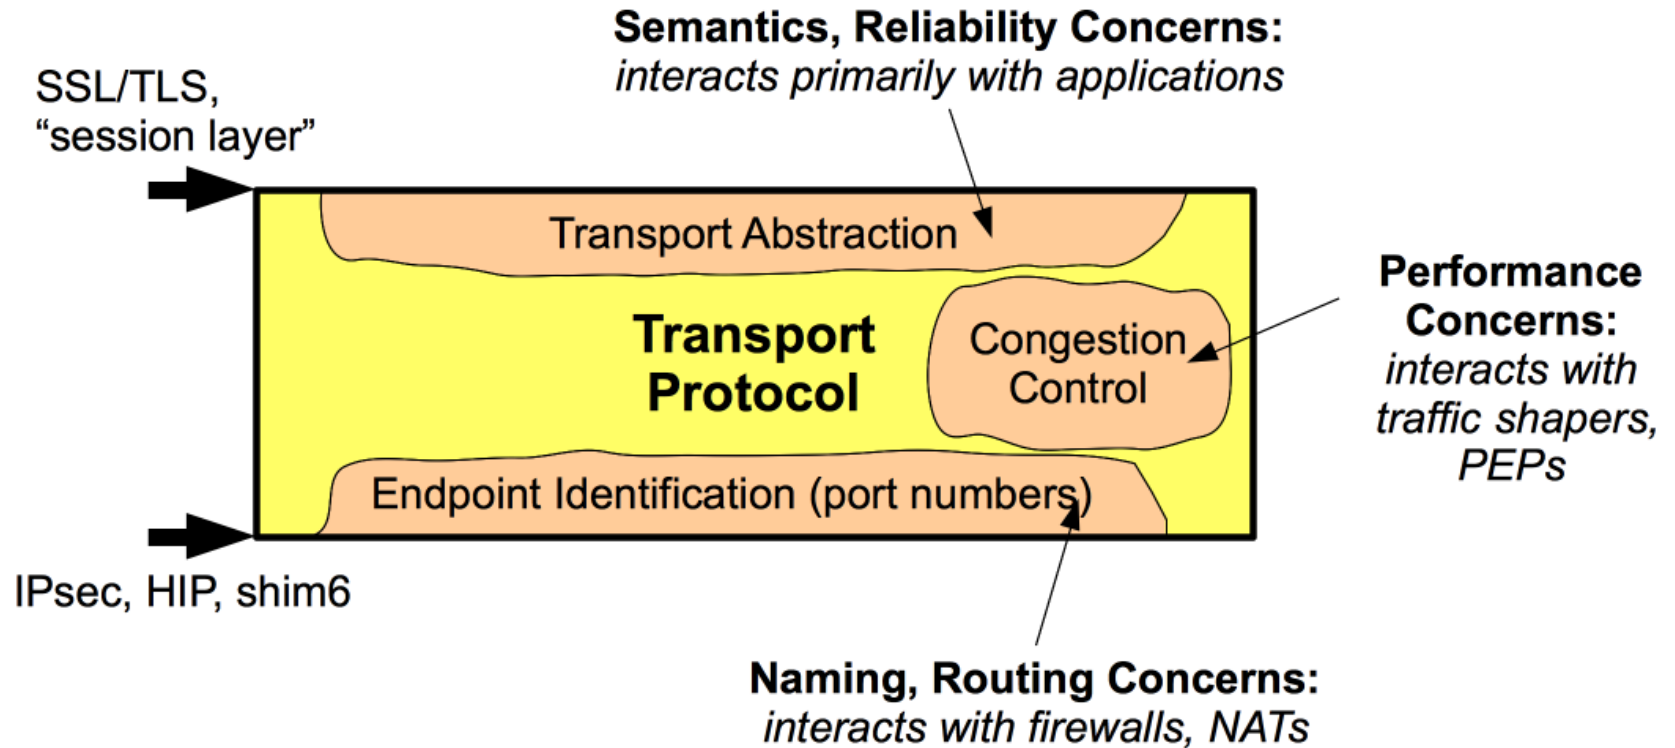
- SCTP: Stream Control Transmission Protocol
- RFC 2960 (circa 2000)
- MPTCP: Multipath with TCP (circa 2012)

Apps build transport abstractions atop TCP, UDP

- Eg: Adobe Flash (RTMFP)

The Transport Layer is Stuck in an Evolutionary Logjam!

The Problem



TCP conflates network-oriented and app-oriented functions

(Thought Experiment:

Design an Internet with middleboxes as first-class citizens)

Design Assumptions For New Transport Services to be Deployable

New end-to-end services should not *require* changes to middleboxes

Consequence: New end-to-end services must appear as legacy protocols on the wire

The Minion Suite

A “packet packhorse” for deploying new transports

- ***Uses legacy protocols ...***

TCP, TLS

- ***... as a substrate...***

turn legacy protocols into *minions* offering unordered datagram service

- ***... for building new services that apps want***

multistreaming, message boundaries, unordered delivery, app-defined congestion control

(may be extended to: stream-level receiver-side flow control, multipath, partial reliability)

The Minion Suite

Main goals:

- How far can we stretch the TCP protocol wire-format?
- TCP increasingly used as a substrate. Make it a good one!

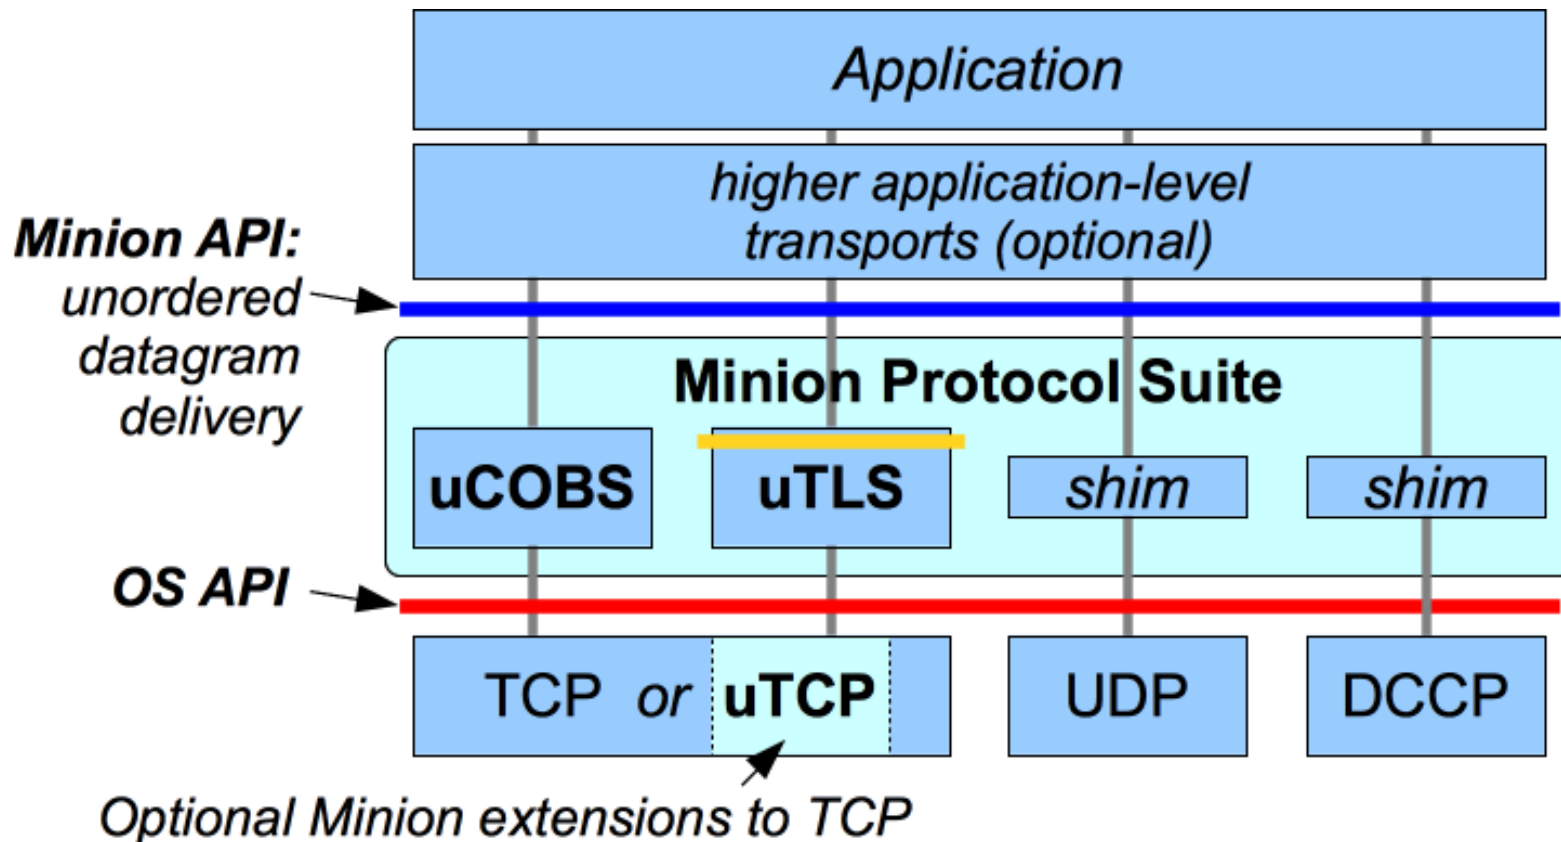
Minion's strength is that it works.

NSDI 2012 review:

A reasonably performing solution to what is (sadly) a real practical problem

[...] accepts the new “narrow waist” as TCP, and shows that with enough devious thought, one can manage to implement functionality at cross-purposes with TCP atop it.

What's in the Minion Suite?



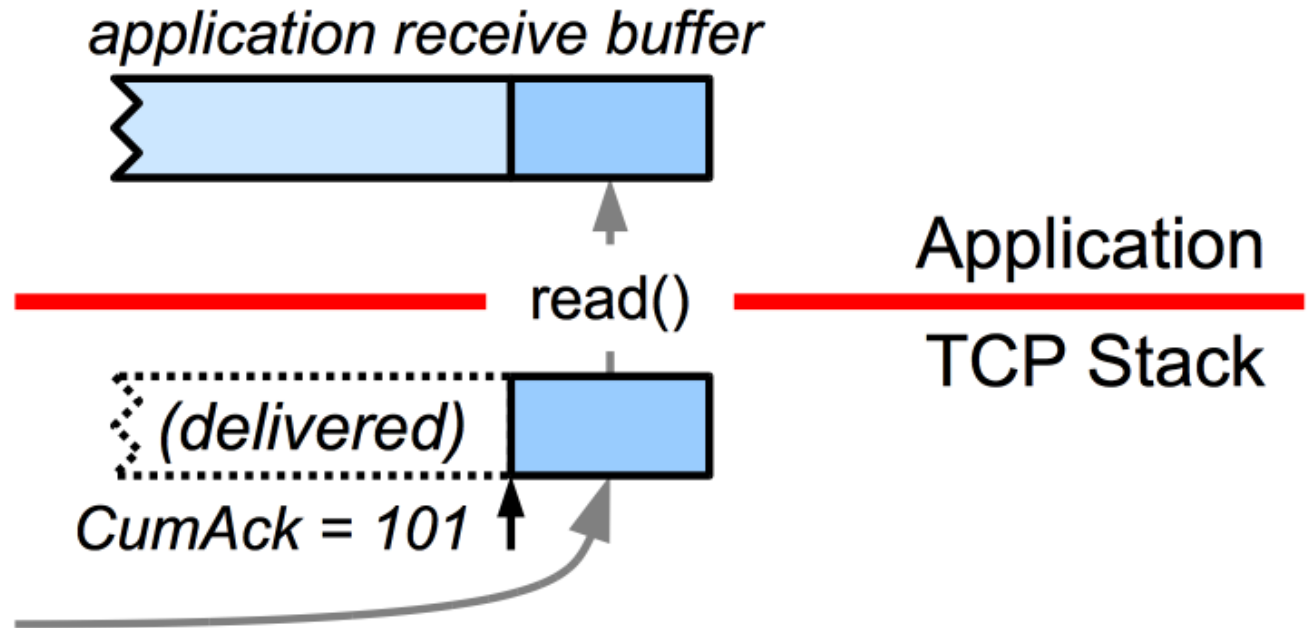

***u*TCP: unordered TCP**

We introduce 2 new TCP socket options in Linux:

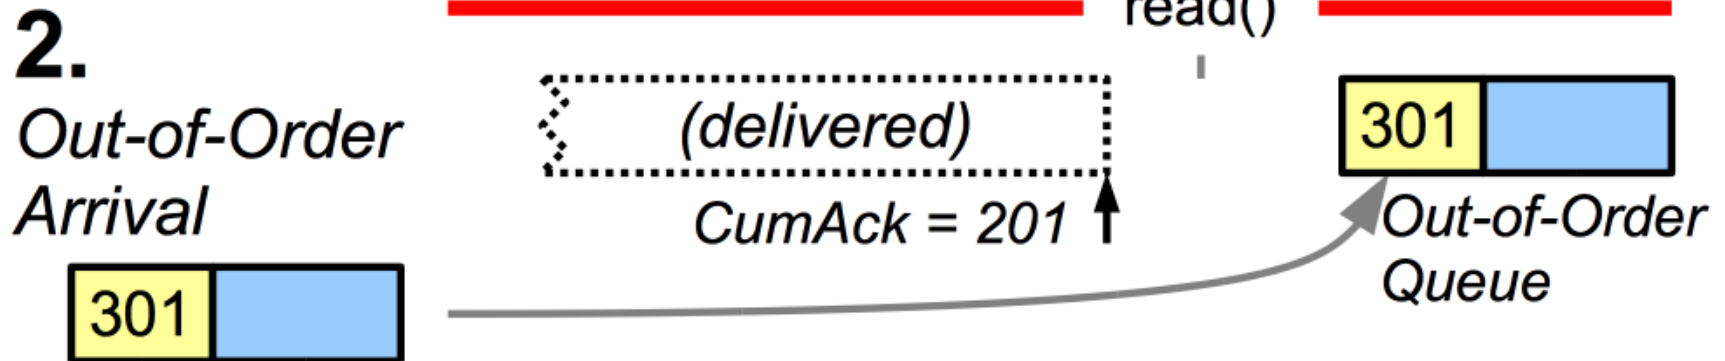
- **SO_UNORDERED_RCV**
 - kernel delivers incoming data immediately
- **SO_UNORDERED_SND**
 - accepts priority with every app message
 - message placed in a priority queue in kernel

Delivery in Standard TCP

1.
*In-Order
Arrival*

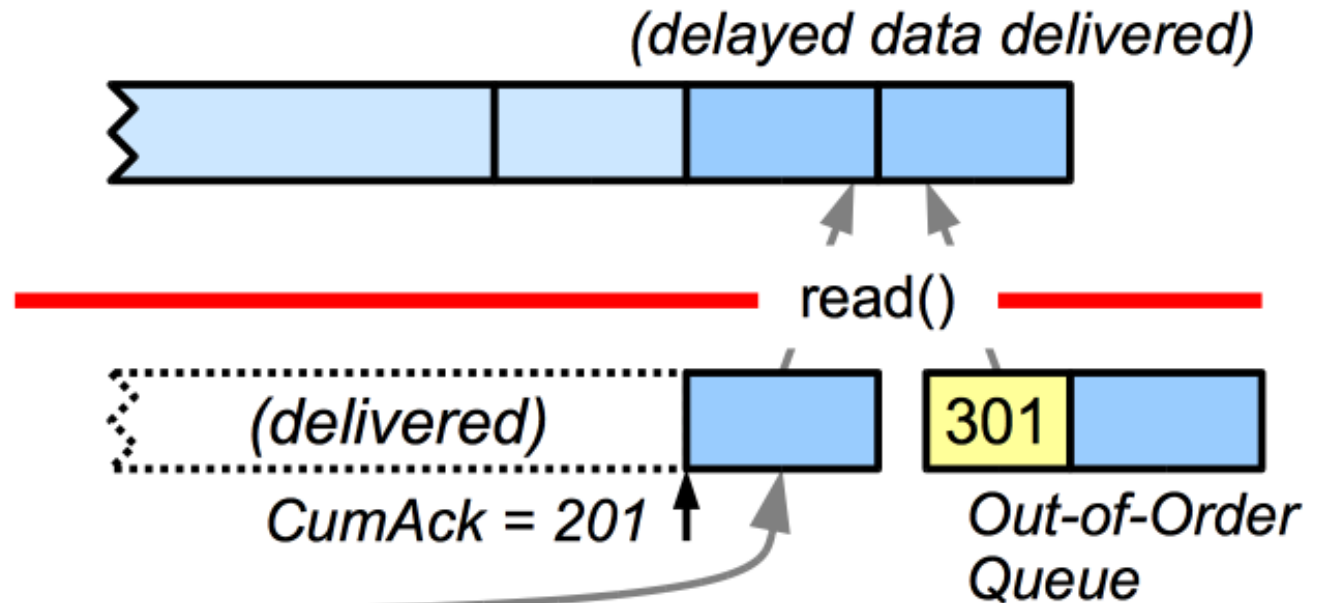



Delivery in Standard TCP

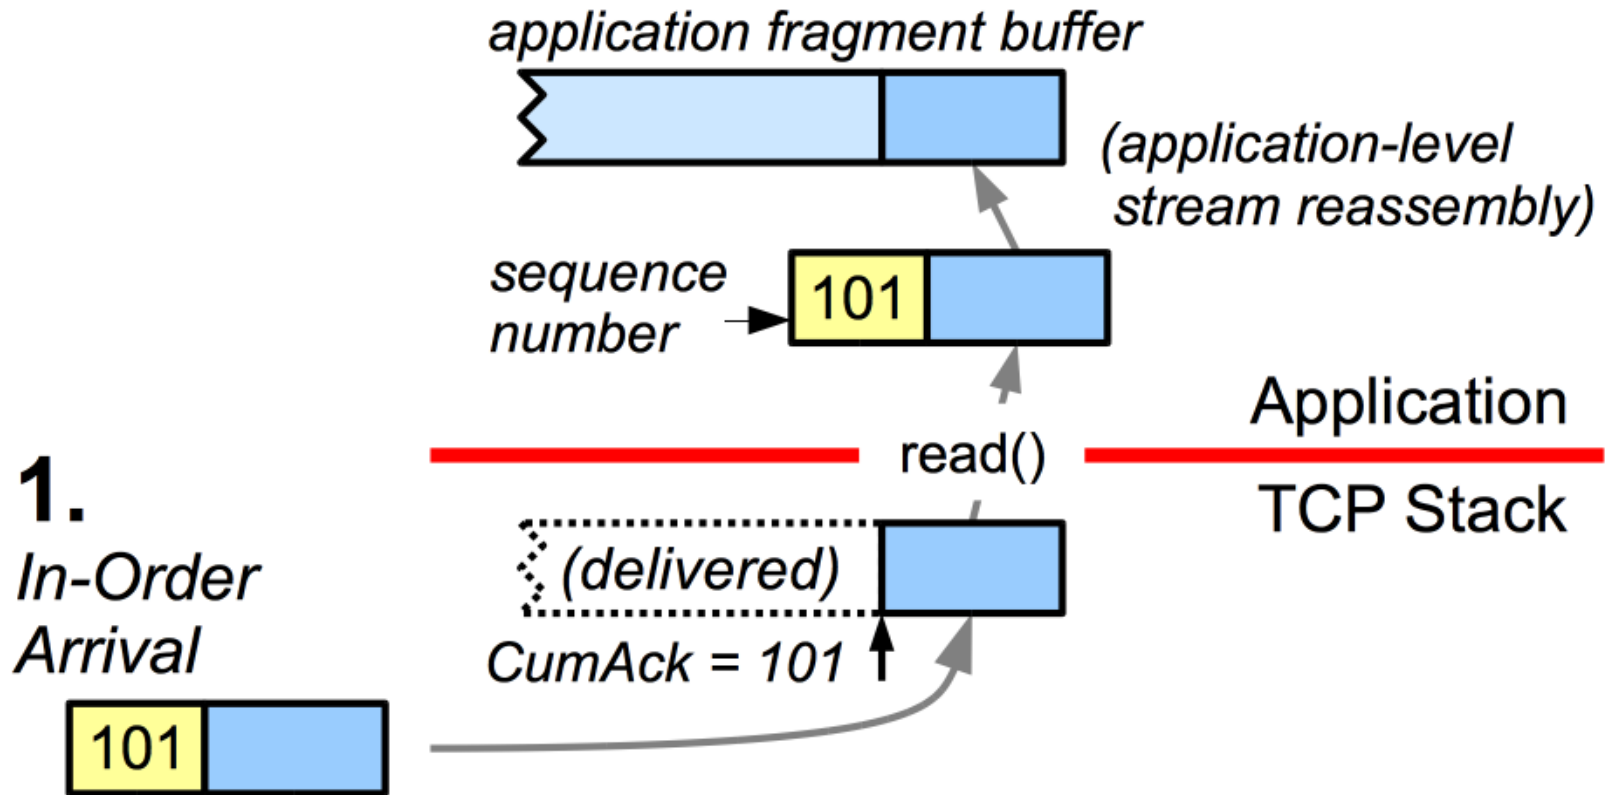


Delivery in Standard TCP

3. Gap-Filling Arrival

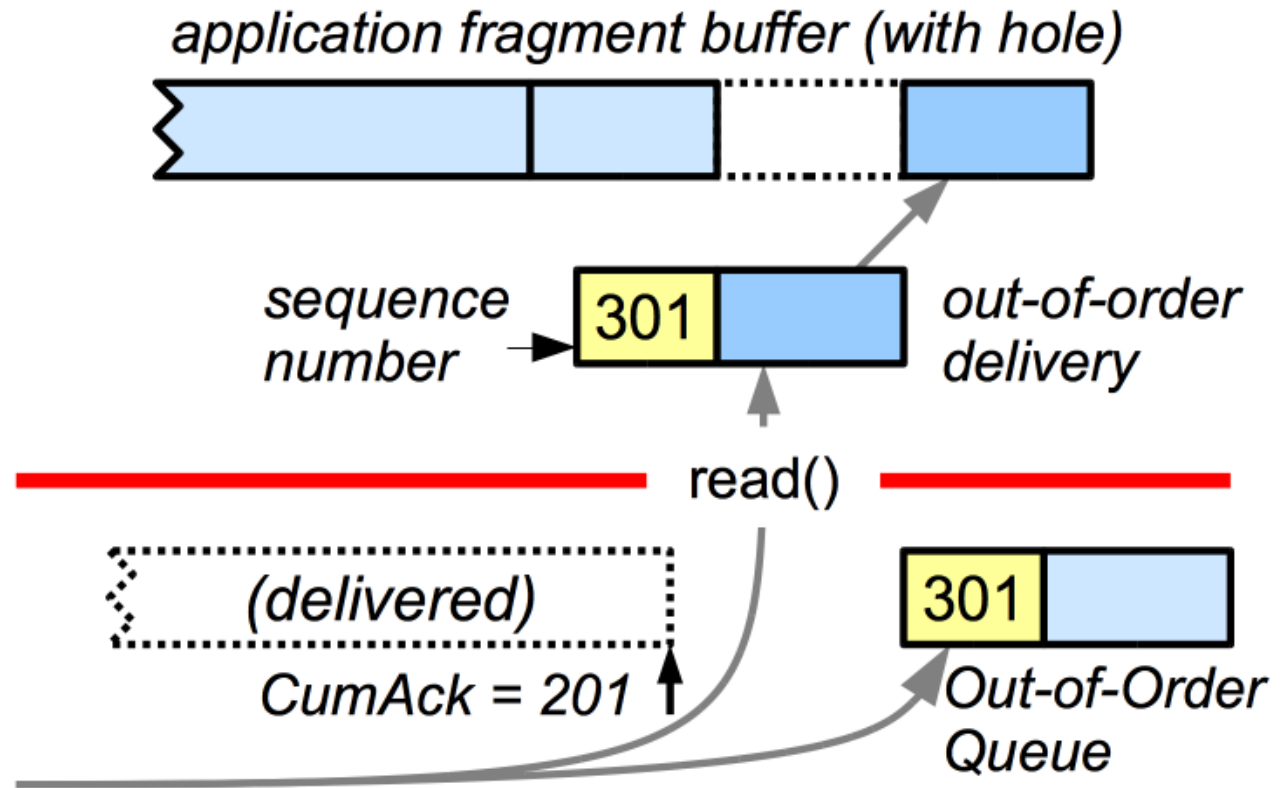



Delivery in *u*TCP



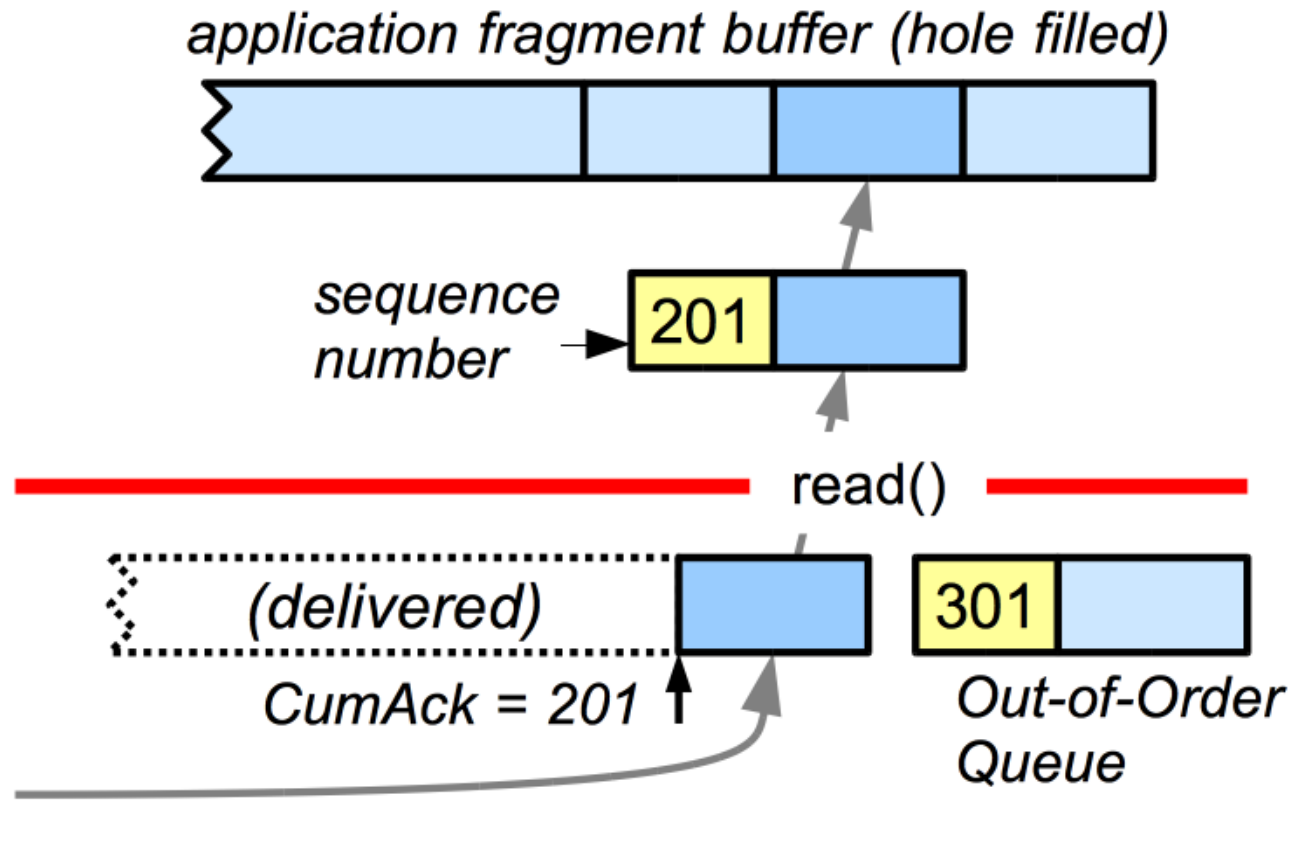

Delivery in *u*TCP

2. Out-of-Order Arrival



Delivery in *u*TCP

3. Gap-Filling Arrival



Dealing with Unordered Bytes

- **No inherent structure in bytestream**
 - *u*TCP receives and delivers arbitrary unordered fragments from the data stream
- ***Self-delimiting* framing with COBS**
 - *zero* added to both ends of an app message
 - COBS encoding eliminates zeros in orig data
 - guaranteed max bit-overhead: 0.4%
(6 bytes for 1448-byte msg)

Using Minion

- ***u*COBS Sender**
 - COBS-encodes messages, sends them with *u*TCP
- ***u*COBS Receiver**
 - manages out-of-order data received from *u*TCP
 - extracts, decodes, delivers messages anywhere in received data bytes


Impact on “Real Applications”

Example: Voice-over-IP (VoIP)

- Delay-sensitive
 - Long RTT delays perceptible, frustrate users
- VoIP codecs are highly sensitive to *burst* losses
 - Can't interpolate when many packets lost or delayed!

VoIP Application: Observed Delay

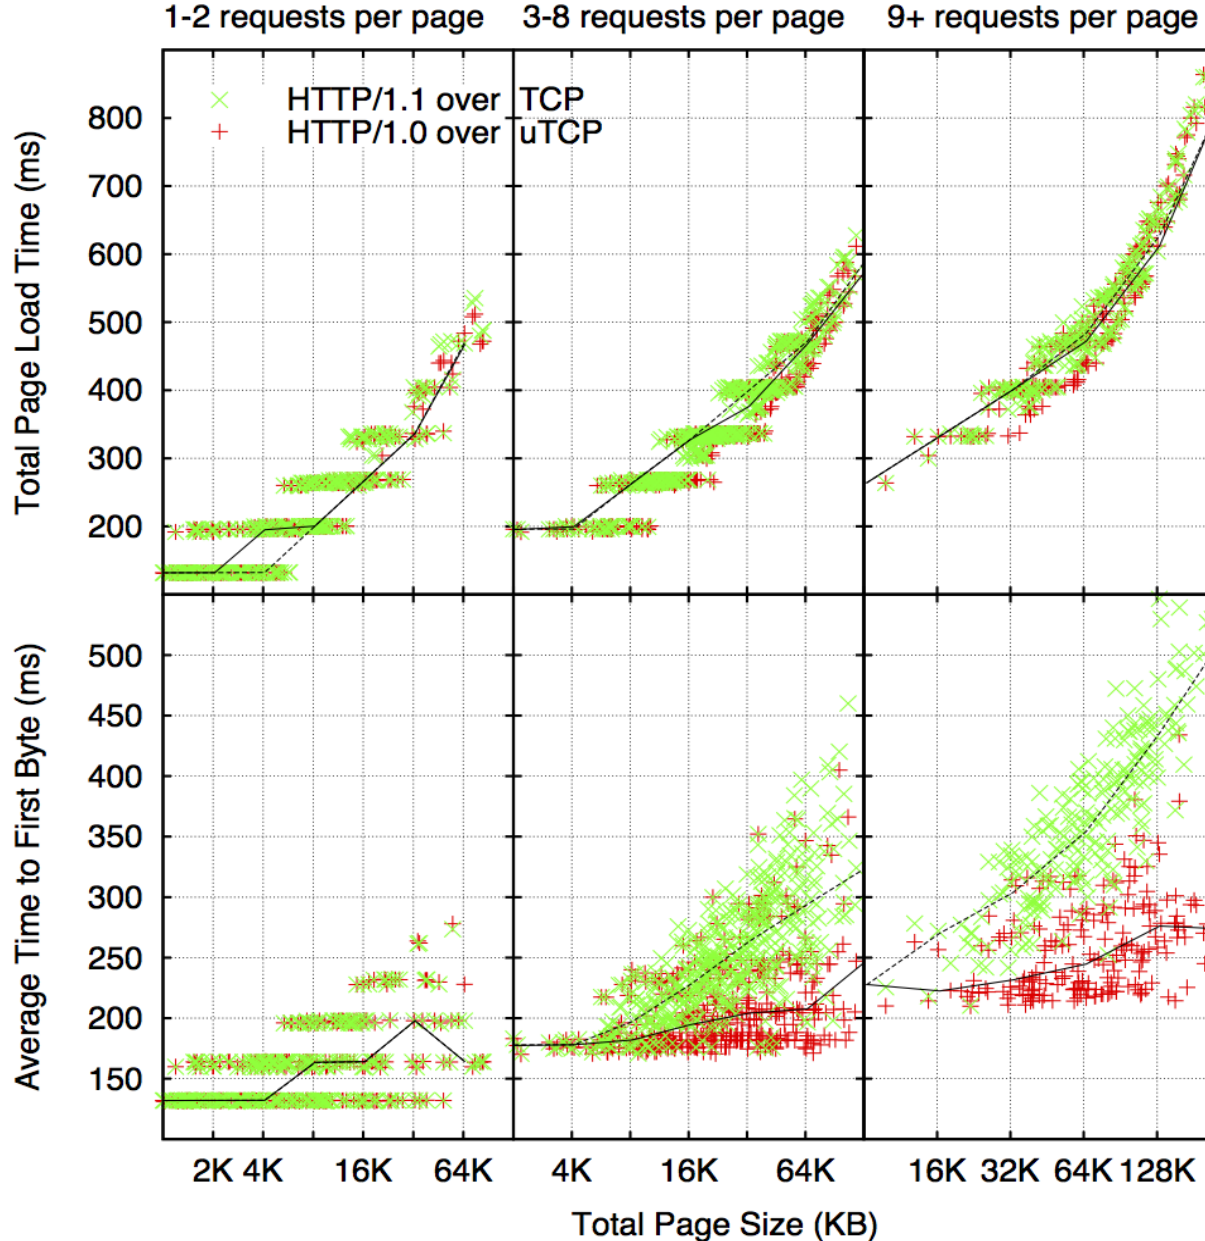
(3Mbps BW; 60ms RTT; 4 TCP flows in background)



Impact on “Real Applications”

- **Example: Web**
 - Independent objects in web pages
 - TCP: parallelism vs. throughput tradeoff
- Multistreaming with Minion
 - ordered streams on top of *u*COBS, 1 per object
 - no Head-of-Line blocking at receiver across streams

User-Perceived Web Latency



Trace-driven,
over a network;
BW of 1.5Mbps,
RTT of 60ms RTT

Much work remains to be done

- Minion goes after transport abstraction issue
 - in an admittedly hack-ish way
 - but aligns incentives well
- Work areas:
 - How to tell reachability between endpoints?
 - Make middleboxes part of the architecture
(Align incentives with incremental deployment)
 - Build a middlebox map of the Internet

Papers, etc.

Look for:

***Tng* (Transport Next Generation), Minion**

<http://www.fandm.edu/jiyengar>

<http://dedis.cs.yale.edu>